

# CSC 311: Introduction to Machine Learning

## Lecture 7 - Probabilistic Models

Murat A. Erdogdu & Richard Zemel

University of Toronto

## Recall: Maximum Likelihood (MLE)

- We have seen in the assignments that various ML algorithms can be derived using the Maximum Likelihood Estimation (MLE).
- Let's start with a simple example: estimating the parameter of a biased coin
  - ▶ You flip a coin  $N = 100$  times. It lands heads  $N_H = 55$  times and tails  $N_T = 45$  times.
  - ▶ What is the probability it will come up heads if we flip again?
- Model: flips are independent Bernoulli random variables with parameter  $\theta$ .
  - ▶ Assume the observations are **independent and identically distributed (i.i.d.)**

# Maximum Likelihood

- The **likelihood function** is the density of the observed data, as a function of parameters  $\theta$ .
- In our case, it's the probability of a *particular* sequence of H/T's.
- Under the Bernoulli model with i.i.d. observations:  
Let  $x_i$  be the # Hs in  $i$ -th flip (can be either 1 or 0)

$$p(x_i = 1|\theta) = \theta \text{ and } p(x_i = 0|\theta) = 1 - \theta$$

$$p(x_i|\theta) = \theta^{x_i} (1 - \theta)^{1-x_i} \quad \text{where } x_i \in \{0, 1\}$$

Likelihood is given as

$$\begin{aligned} L(\theta) &= p(x_1, \dots, x_N|\theta) = \prod_{i=1}^N \theta^{x_i} (1 - \theta)^{1-x_i} \\ &= \theta^{N_H} (1 - \theta)^{N_T} \end{aligned}$$

where  $N_H = \sum_i x_i$  and  $N_T = N - \sum_i x_i$

- We usually work with log-likelihoods:

$$\ell(\theta) = \log L(\theta) = N_H \log \theta + N_T \log(1 - \theta)$$

# Maximum Likelihood

- Good values of  $\theta$  should assign high probability to the observed data. This motivates the **maximum likelihood criterion**.
- Remember how we found the optimal solution to linear regression by setting derivatives to zero? We can do that again for the coin example.

$$\begin{aligned}\frac{d\ell}{d\theta} &= \frac{d}{d\theta} (N_H \log \theta + N_T \log(1 - \theta)) \\ &= \frac{N_H}{\theta} - \frac{N_T}{1 - \theta}\end{aligned}$$

- Setting this to zero gives the maximum likelihood estimate:

$$\hat{\theta}_{\text{ML}} = \frac{N_H}{N_H + N_T}.$$

# Generative vs Discriminative

Two approaches to classification:

- **Discriminative approach:** estimate parameters of decision boundary/class separator directly from labeled examples.
  - ▶ Tries to solve: How do I separate the classes?
  - ▶ learn  $p(y|x)$  directly (logistic regression models)
  - ▶ learn mappings from inputs to classes (linear/logistic regression, decision trees etc)
- **Generative approach:** model the distribution of inputs characteristic of the class (Bayes classifier).
  - ▶ Tries to solve: What does each class "look" like?
  - ▶ Build a model of  $p(x|y)$
  - ▶ Apply Bayes Rule
- Key difference: is there a distributional assumption over inputs?

# A Generative Model: Bayes Classifier

- Aim to classify text into spam/not-spam (yes  $c=1$ ; no  $c=0$ )
- Example: “You are one of the very few who have been selected as a winners for the free \$1000 Gift Card.”
- Use bag-of-words features, get binary vector  $\mathbf{x}$  for each email
- Vocabulary:
  - ▶ “a”: 1
  - ▶ ...
  - ▶ “car”: 0
  - ▶ “card”: 1
  - ▶ ...
  - ▶ “win”: 0
  - ▶ “winner”: 1
  - ▶ “winter”: 0
  - ▶ ...
  - ▶ “you”: 1

# Bayes Classifier

- Given features  $\mathbf{x} = [x_1, x_2, \dots, x_D]^T$  we want to compute class probabilities using Bayes Rule:

$$\underbrace{p(c|\mathbf{x})}_{\text{Pr. class given words}} = \frac{p(\mathbf{x}, c)}{p(\mathbf{x})} = \frac{\overbrace{p(\mathbf{x}|c)}^{\text{Pr. words given class}}}{\overbrace{p(\mathbf{x})}^{\text{Pr.}}} p(c)$$

- More formally

$$\text{posterior} = \frac{\text{Class likelihood} \times \text{prior}}{\text{Evidence}}$$

- How can we compute  $p(\mathbf{x})$  for the two class case? (Do we need to?)

$$p(\mathbf{x}) = p(\mathbf{x}|c=0)p(c=0) + p(\mathbf{x}|c=1)p(c=1)$$

- To compute  $p(c|\mathbf{x})$  we need:  $p(\mathbf{x}|c)$  and  $p(c)$

# Naïve Bayes

- Assume we have two classes: spam and non-spam. We have a dictionary of  $D$  words, and binary features  $\mathbf{x} = [x_1, \dots, x_D]$  saying whether each word appears in the e-mail.
- If we define a joint distribution  $p(c, x_1, \dots, x_D)$ , this gives enough information to determine  $p(c)$  and  $p(\mathbf{x}|c)$ .
- Problem: specifying a joint distribution over  $D + 1$  binary variables requires  $2^{D+1} - 1$  entries. This is computationally prohibitive and would require an absurd amount of data to fit.
- We'd like to impose **structure** on the distribution such that:
  - ▶ it can be **compactly** represented
  - ▶ **learning** and **inference** are both tractable

# Naïve Bayes


- Naïve assumption: **Naïve Bayes** assumes that the word features  $x_i$  are **conditionally independent** given the class  $c$ .
  - ▶ This means  $x_i$  and  $x_j$  are independent under the conditional distribution  $p(\mathbf{x}|c)$ .
  - ▶ Note: this doesn't mean they're independent.
  - ▶ Mathematically,

$$p(c, x_1, \dots, x_D) = p(c)p(x_1|c) \cdots p(x_D|c).$$

- Compact representation of the joint distribution
  - ▶ Prior probability of class:  $p(c = 1) = \pi$  (e.g. spam email)
  - ▶ Conditional probability of word feature given class:  
 $p(x_j = 1|c) = \theta_{jc}$  (e.g. word "price" appearing in spam)
  - ▶  $2D + 1$  parameters total (before  $2^{D+1} - 1$ )

# Bayes Nets

- We can represent this model using an **directed graphical model**, or **Bayesian network**:



- This graph structure means the joint distribution factorizes as a product of conditional distributions for each variable given its parent(s).
- Intuitively, you can think of the edges as reflecting a causal structure. But mathematically, this doesn't hold without additional assumptions.

# Naïve Bayes: Learning

- The parameters can be learned efficiently because the log-likelihood decomposes into independent terms for each feature.

$$\begin{aligned}\ell(\boldsymbol{\theta}) &= \sum_{i=1}^N \log p(c^{(i)}, \mathbf{x}^{(i)}) = \sum_{i=1}^N \log \left\{ p(\mathbf{x}^{(i)} | c^{(i)}) p(c^{(i)}) \right\} \\ &= \sum_{i=1}^N \log \left\{ p(c^{(i)}) \prod_{j=1}^D p(x_j^{(i)} | c^{(i)}) \right\} \\ &= \sum_{i=1}^N \left[ \log p(c^{(i)}) + \sum_{j=1}^D \log p(x_j^{(i)} | c^{(i)}) \right] \\ &= \underbrace{\sum_{i=1}^N \log p(c^{(i)})}_{\text{Bernoulli log-likelihood of labels}} + \underbrace{\sum_{j=1}^D \sum_{i=1}^N \log p(x_j^{(i)} | c^{(i)})}_{\text{Bernoulli log-likelihood for feature } x_j}\end{aligned}$$

- Each of these log-likelihood terms depends on different sets of parameters, so they can be optimized independently.

## Naïve Bayes: Learning

- We can handle these terms separately. For the prior we maximize:  
 $\sum_{i=1}^N \log p(c^{(i)})$
- This is a minor variant of our coin flip example. Let  $p(c^{(i)} = 1) = \pi$ . Note  $p(c^{(i)}) = \pi^{c^{(i)}} (1 - \pi)^{1 - c^{(i)}}$ .
- Log-likelihood:

$$\sum_{i=1}^N \log p(c^{(i)}) = \sum_{i=1}^N c^{(i)} \log \pi + \sum_{i=1}^N (1 - c^{(i)}) \log(1 - \pi)$$

- Obtain MLEs by setting derivatives to zero:

$$\pi = \frac{\sum_i \mathbb{I}[c^{(i)} = 1]}{N} = \frac{\text{\# spams in dataset}}{\text{total \# samples}}$$

# Naïve Bayes: Learning

- Each  $\theta_{jc}$ 's can be treated separately: maximize  $\sum_{i=1}^N \log p(x_j^{(i)} | c^{(i)})$
- This is (again) a minor variant of our coin flip example.

Let  $\theta_{jc} = p(x_j^{(i)} = 1 | c)$ . Note  $p(x_j^{(i)} | c) = \theta_{jc}^{x_j^{(i)}} (1 - \theta_{jc})^{1-x_j^{(i)}}$ .

- Log-likelihood:

$$\begin{aligned}\sum_{i=1}^N \log p(x_j^{(i)} | c^{(i)}) &= \sum_{i=1}^N c^{(i)} \left\{ x_j^{(i)} \log \theta_{j1} + (1 - x_j^{(i)}) \log (1 - \theta_{j1}) \right\} \\ &\quad + \sum_{i=1}^N (1 - c^{(i)}) \left\{ x_j^{(i)} \log \theta_{j0} + (1 - x_j^{(i)}) \log (1 - \theta_{j0}) \right\}\end{aligned}$$

- Obtain MLEs by setting derivatives to zero:

$$\theta_{jc} = \frac{\sum_i \mathbb{I}[x_j^{(i)} = 1 \text{ } \& \text{ } c^{(i)} = c]}{\sum_i \mathbb{I}[c^{(i)} = c]} = \frac{\#\text{word } j \text{ appears in spams}}{\# \text{ spams in dataset}}$$

## Naïve Bayes: Inference

- We predict the category by performing **inference** in the model.
- Apply **Bayes' Rule**:

$$p(c \mid \mathbf{x}) = \frac{p(c)p(\mathbf{x} \mid c)}{\sum_{c'} p(c')p(\mathbf{x} \mid c')} = \frac{p(c) \prod_{j=1}^D p(x_j \mid c)}{\sum_{c'} p(c') \prod_{j=1}^D p(x_j \mid c')}$$

- We need not compute the denominator if we're simply trying to determine the most likely  $c$ .
- Shorthand notation:

$$p(c \mid \mathbf{x}) \propto p(c) \prod_{j=1}^D p(x_j \mid c)$$

- For input  $\mathbf{x}$ , predict by comparing the values of  $p(c) \prod_{j=1}^D p(x_j \mid c)$  for different  $c$  (e.g. choose the largest).

# Naïve Bayes

- Naïve Bayes is an amazingly cheap learning algorithm!
- **Training time:** estimate parameters using maximum likelihood
  - ▶ Compute co-occurrence counts of each feature with the labels.
  - ▶ Requires only one pass through the data!
- **Test time:** apply Bayes' Rule
  - ▶ Cheap because of the model structure. (For more general models, Bayesian inference can be very expensive and/or complicated.)
- We covered the Bernoulli case for simplicity. But our analysis easily extends to other probability distributions.
- Unfortunately, it's usually less accurate in practice compared to discriminative models due to its “naïve” independence assumption.

## MLE issue: Data Sparsity


- Maximum likelihood has a pitfall: if you have too little data, it can overfit.
- E.g., what if you flip the coin twice and get H both times?

$$\theta_{\text{ML}} = \frac{N_H}{N_H + N_T} = \frac{2}{2 + 0} = 1$$


- Because it never observed T, it assigns this outcome probability 0. This problem is known as **data sparsity**.

# Bayesian Parameter Estimation

- In maximum likelihood, the observations are treated as random variables, but the parameters are not.



- The **Bayesian** approach treats the parameters as random variables as well.  $\beta$  is the set of parameters in the prior distribution of  $\theta$ .



- To define a Bayesian model, we need to specify two distributions:
  - ▶ The **prior distribution**  $p(\theta)$ , which encodes our beliefs about the parameters *before* we observe the data
  - ▶ The **likelihood**  $p(\mathcal{D} | \theta)$ , same as in maximum likelihood

# Bayesian Parameter Estimation

- When we **update** our beliefs based on the observations, we compute the **posterior distribution** using Bayes' Rule:

$$p(\boldsymbol{\theta} | \mathcal{D}) = \frac{p(\boldsymbol{\theta})p(\mathcal{D} | \boldsymbol{\theta})}{\int p(\boldsymbol{\theta}')p(\mathcal{D} | \boldsymbol{\theta}') d\boldsymbol{\theta}'}.$$

- We rarely ever compute the denominator explicitly. In general, it is computationally intractable.

# Bayesian Parameter Estimation

- Let's revisit the coin example. We already know the likelihood:

$$L(\theta) = p(\mathcal{D}|\theta) = \theta^{N_H} (1 - \theta)^{N_T}$$

- It remains to specify the prior  $p(\theta)$ .
  - We can choose an **uninformative prior**, which assumes as little as possible. A reasonable choice is the uniform prior.
  - But our experience tells us 0.5 is more likely than 0.99. One particularly useful prior that lets us specify this is the **beta distribution**:


$$p(\theta; a, b) = \frac{\Gamma(a + b)}{\Gamma(a)\Gamma(b)} \theta^{a-1} (1 - \theta)^{b-1}.$$

- This notation for proportionality lets us ignore the normalization constant:

$$p(\theta; a, b) \propto \theta^{a-1} (1 - \theta)^{b-1}.$$

# Bayesian Parameter Estimation

- Beta distribution for various values of  $a, b$ :



- Some observations:
  - ▶ The expectation  $\mathbb{E}[\theta] = a/(a + b)$  (easy to derive).
  - ▶ The distribution gets more peaked when  $a$  and  $b$  are large.
  - ▶ The uniform distribution is the special case where  $a = b = 1$ .
- The beta distribution is used for is as a prior for the Bernoulli distribution.

# Bayesian Parameter Estimation

- Computing the posterior distribution:

$$\begin{aligned} p(\boldsymbol{\theta} | \mathcal{D}) &\propto p(\boldsymbol{\theta})p(\mathcal{D} | \boldsymbol{\theta}) \\ &\propto \left[\theta^{a-1}(1-\theta)^{b-1}\right] \left[\theta^{N_H}(1-\theta)^{N_T}\right] \\ &= \theta^{a-1+N_H}(1-\theta)^{b-1+N_T}. \end{aligned}$$

- This is just a beta distribution with parameters  $N_H + a$  and  $N_T + b$ .
- The posterior expectation of  $\theta$  is:

$$\mathbb{E}[\theta | \mathcal{D}] = \frac{N_H + a}{N_H + N_T + a + b}$$


- The parameters  $a$  and  $b$  of the prior can be thought of as **pseudo-counts**.
  - ▶ The reason this works is that the prior and likelihood have the same functional form. This phenomenon is known as **conjugacy** (conjugate priors), and it's very useful.

# Bayesian Parameter Estimation

Bayesian inference for the coin flip example:


Small data setting

$$N_H = 2, N_T = 0$$



Large data setting


$$N_H = 55, N_T = 45$$



When you have enough observations, the **data overwhelm the prior**.

# Maximum A-Posteriori Estimation

- **Maximum a-posteriori (MAP) estimation:** find the most likely parameter settings under the posterior



# Maximum A-Posteriori Estimation

- This converts the Bayesian parameter estimation problem into a maximization problem

$$\begin{aligned}\hat{\boldsymbol{\theta}}_{\text{MAP}} &= \arg \max_{\boldsymbol{\theta}} p(\boldsymbol{\theta} \mid \mathcal{D}) \\&= \arg \max_{\boldsymbol{\theta}} p(\boldsymbol{\theta}, \mathcal{D}) \\&= \arg \max_{\boldsymbol{\theta}} p(\boldsymbol{\theta}) p(\mathcal{D} \mid \boldsymbol{\theta}) \\&= \arg \max_{\boldsymbol{\theta}} \log p(\boldsymbol{\theta}) + \log p(\mathcal{D} \mid \boldsymbol{\theta})\end{aligned}$$

- We already saw an example of this in the homework.

# Maximum A-Posteriori Estimation

- Joint probability in the coin flip example:

$$\begin{aligned}\log p(\theta, \mathcal{D}) &= \log p(\theta) + \log p(\mathcal{D} | \theta) \\ &= \text{Const} + (a - 1) \log \theta + (b - 1) \log(1 - \theta) + N_H \log \theta + N_T \log(1 - \theta) \\ &= \text{Const} + (N_H + a - 1) \log \theta + (N_T + b - 1) \log(1 - \theta)\end{aligned}$$

- Maximize by finding a critical point

$$0 = \frac{d}{d\theta} \log p(\theta, \mathcal{D}) = \frac{N_H + a - 1}{\theta} - \frac{N_T + b - 1}{1 - \theta}$$

- Solving for  $\theta$ ,

$$\hat{\theta}_{\text{MAP}} = \frac{N_H + a - 1}{N_H + N_T + a + b - 2}$$

# Maximum A-Posteriori Estimation

Comparison of estimates in the coin flip example:

Formula	$N_H = 2, N_T = 0$	$N_H = 55, N_T = 45$
$\hat{\theta}_{\text{ML}}$	$\frac{N_H}{N_H + N_T}$	1 $\frac{55}{100} = 0.55$
$\mathbb{E}[\theta   \mathcal{D}]$	$\frac{N_H + a}{N_H + N_T + a + b}$	$\frac{4}{6} \approx 0.67$ $\frac{57}{104} \approx 0.548$
$\hat{\theta}_{\text{MAP}}$	$\frac{N_H + a - 1}{N_H + N_T + a + b - 2}$	$\frac{3}{4} = 0.75$ $\frac{56}{102} \approx 0.549$

$\hat{\theta}_{\text{MAP}}$  assigns nonzero probabilities as long as  $a, b > 1$ .

# Gaussian Discriminant Analysis

- Generative models - data generating distribution  $p(\mathbf{x}|t = k)$
- Instead of trying to separate classes, try to model what each class "looks like".
- Recall that  $p(\mathbf{x}|t = k)$  may be very complex

$$p(x_1, \dots, x_d, y) = p(x_1|x_2, \dots, x_d, y) \cdots p(x_{d-1}|x_d, y)p(x_d, y)$$

- Naive bayes used a conditional independence assumption. What else could we do? Choose a simple distribution.
- Next, we will discuss fitting Gaussian distributions to our data.

# Bayes Classifier


- Let's take a step back...
- Bayes Classifier

$$\begin{aligned} h(\mathbf{x}) &= \arg \max_k p(t = k | \mathbf{x}) = \arg \max \frac{p(\mathbf{x}|t = k)p(t = k)}{p(\mathbf{x})} \\ &= \arg \max_k p(\mathbf{x}|t = k)p(t = k) \end{aligned}$$

- Talked about Discrete  $\mathbf{x}$ , what if  $\mathbf{x}$  is continuous?

# Classification: Diabetes Example

- Observation per patient: White blood cell count & glucose value.



- How can we model  $p(x|t = k)$ ? Multivariate Gaussian

# Multivariate Data

- Multiple measurements (sensors)
- $D$  inputs/features/attributes
- $N$  instances/observations/examples

$$\mathbf{X} = \begin{bmatrix} [\mathbf{x}^{(1)}]^\top \\ [\mathbf{x}^{(2)}]^\top \\ \vdots \\ [\mathbf{x}^{(N)}]^\top \end{bmatrix} = \begin{bmatrix} x_1^{(1)} & x_2^{(1)} & \cdots & x_D^{(1)} \\ x_1^{(2)} & x_2^{(2)} & \cdots & x_D^{(2)} \\ \vdots & \vdots & \ddots & \vdots \\ x_1^{(N)} & x_2^{(N)} & \cdots & x_D^{(N)} \end{bmatrix}$$

# Multivariate Parameters

- Mean

$$\mathbb{E}[\mathbf{x}^{(i)}] = \boldsymbol{\mu} = [\mu_1, \dots, \mu_d]^T \in \mathbb{R}^D$$

- Covariance


$$\boldsymbol{\Sigma} = Cov(\mathbf{x}^{(i)}) = \mathbb{E}[(\mathbf{x}^{(i)} - \boldsymbol{\mu})(\mathbf{x}^{(i)} - \boldsymbol{\mu})^\top] = \begin{bmatrix} \sigma_1^2 & \sigma_{12} & \cdots & \sigma_{1D} \\ \sigma_{12} & \sigma_2^2 & \cdots & \sigma_{2D} \\ \vdots & \vdots & \ddots & \vdots \\ \sigma_{D1} & \sigma_{D2} & \cdots & \sigma_D^2 \end{bmatrix}$$

- For Gaussians - all you need to know to represent (not true in general).

# Multivariate Gaussian Distribution

- $\mathbf{x} \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ , a Gaussian (or normal) distribution defined as

$$p(\mathbf{x}) = \frac{1}{(2\pi)^{d/2} |\boldsymbol{\Sigma}|^{1/2}} \exp \left[ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right]$$



- The Central Limit Theorem says that sums of lots of independent random variables are approximately Gaussian.
- In machine learning, we use Gaussians a lot because they make the calculations easy.

# Bivariate Normal

$$\Sigma = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$$

$$\Sigma = 0.5 \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$$

$$\Sigma = 2 \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$$




Figure: Probability density function




Figure: Contour plot of the pdf

# Bivariate Normal

$$\Sigma = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$$

$$\Sigma = \begin{pmatrix} 1 & 0.5 \\ 0.5 & 1 \end{pmatrix}$$

$$\Sigma = \begin{pmatrix} 1 & 0.8 \\ 0.8 & 1 \end{pmatrix}$$




Figure: Probability density function




Figure: Contour plot of the pdf

# Maximum Likelihood

- Suppose we want to model the distribution of highest and lowest temperatures in Toronto in March, and we've recorded the following observations :  
 $(-2.5, -7.5) \quad (-9.9, -14.9) \quad (-12.1, -17.5) \quad (-8.9, -13.9) \quad (-6.0, -11.1)$
- Assume they're drawn from a Gaussian distribution with mean  $\mu$ , and covariance  $\Sigma$ . We want to estimate these using data.
- Log-likelihood function:

$$\begin{aligned}\ell(\mu, \Sigma) &= \log \prod_{i=1}^N \left[ \frac{1}{(2\pi)^{d/2} |\Sigma|^{1/2}} \exp \left\{ -\frac{1}{2} (\mathbf{x}^{(i)} - \mu)^T \Sigma^{-1} (\mathbf{x}^{(i)} - \mu) \right\} \right] \\ &= \sum_{i=1}^N \log \left[ \frac{1}{(2\pi)^{d/2} |\Sigma|^{1/2}} \exp \left\{ -\frac{1}{2} (\mathbf{x}^{(i)} - \mu)^T \Sigma^{-1} (\mathbf{x}^{(i)} - \mu) \right\} \right] \\ &= \sum_{i=1}^N \underbrace{-\log(2\pi)^{d/2} - \log |\Sigma|^{1/2}}_{\text{constant}} - \frac{1}{2} (\mathbf{x}^{(i)} - \mu)^T \Sigma^{-1} (\mathbf{x}^{(i)} - \mu)\end{aligned}$$

# Maximum Likelihood

- Maximize the log-likelihood by setting the derivative to zero:

$$\begin{aligned} 0 &= \frac{d\ell}{d\boldsymbol{\mu}} = - \sum_{i=1}^N \frac{d}{d\boldsymbol{\mu}} \frac{1}{2} (\mathbf{x}^{(i)} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x}^{(i)} - \boldsymbol{\mu}) \\ &= - \sum_{i=1}^N \boldsymbol{\Sigma}^{-1} (\mathbf{x}^{(i)} - \boldsymbol{\mu}) = 0 \end{aligned}$$

- Solving we get  $\hat{\boldsymbol{\mu}} = \frac{1}{N} \sum_{i=1}^N \mathbf{x}^{(i)}$ . In general, “hat” means estimator
- This is just the sample mean of the observed values, or the **empirical mean**.

# Maximum Likelihood

- Similar calculation for the covariance matrix  $\Sigma$  yields:
- Set the *partial* derivatives to zero, just like before

$$0 = \frac{\partial \ell}{\partial \Sigma} \implies \hat{\Sigma} = \frac{1}{N} \sum_{i=1}^N (\mathbf{x}^{(i)} - \hat{\mu})(\mathbf{x}^{(i)} - \hat{\mu})^\top$$

- This is called the empirical covariance and comes up quite often (i.e. PCA next lecture)
- Derivation in multivariate case is tedious. No need to worry about it. But it is good practice to derive this in one dimension. See appendix.

# Gaussian Discriminant Analysis (Gaussian Bayes Classifier)

- Gaussian Discriminant Analysis in its general form assumes that  $p(\mathbf{x}|t)$  is distributed according to a multivariate normal (Gaussian) distribution
- Multivariate Gaussian distribution:

$$p(\mathbf{x}|t = k) = \frac{1}{(2\pi)^{d/2} |\boldsymbol{\Sigma}_k|^{1/2}} \exp \left[ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu}_k)^T \boldsymbol{\Sigma}_k^{-1} (\mathbf{x} - \boldsymbol{\mu}_k) \right]$$

where  $|\boldsymbol{\Sigma}_k|$  denotes the determinant of the matrix, and  $D$  is dimension of  $\mathbf{x}$

- Each class  $k$  has a mean vector  $\boldsymbol{\mu}_k$  and a covariance matrix  $\boldsymbol{\Sigma}_k$
- $\boldsymbol{\Sigma}_k$  has  $\mathcal{O}(D^2)$  parameters - could be hard to estimate

# Gaussian Discriminant Analysis (Gaussian Bayes Classifier)

- GDA (GBC) decision boundary is based on class posterior.
- Make decisions by comparing class probabilities:


$$\begin{aligned}\log p(t_k|\mathbf{x}) &= \log p(\mathbf{x}|t_k) + \log p(t_k) - \log p(\mathbf{x}) \\ &= -\frac{d}{2} \log(2\pi) - \frac{1}{2} \log |\boldsymbol{\Sigma}_k^{-1}| - \frac{1}{2} (\mathbf{x} - \boldsymbol{\mu}_k)^T \boldsymbol{\Sigma}_k^{-1} (\mathbf{x} - \boldsymbol{\mu}_k) \\ &\quad + \log p(t_k) - \log p(\mathbf{x})\end{aligned}$$

- Decision boundary ( $\log p(t_k|\mathbf{x}) = \log p(t_l|\mathbf{x})$ ):

$$\begin{aligned}(\mathbf{x} - \boldsymbol{\mu}_k)^T \boldsymbol{\Sigma}_k^{-1} (\mathbf{x} - \boldsymbol{\mu}_k) &= (\mathbf{x} - \boldsymbol{\mu}_l)^T \boldsymbol{\Sigma}_l^{-1} (\mathbf{x} - \boldsymbol{\mu}_l) + C_{k,l} \\ \mathbf{x}^T \boldsymbol{\Sigma}_k^{-1} \mathbf{x} - 2\boldsymbol{\mu}_k^T \boldsymbol{\Sigma}_k^{-1} \mathbf{x} &= \mathbf{x}^T \boldsymbol{\Sigma}_l^{-1} \mathbf{x} - 2\boldsymbol{\mu}_l^T \boldsymbol{\Sigma}_l^{-1} \mathbf{x} + C_{k,l}\end{aligned}$$

- Quadratic function in  $\mathbf{x} \implies$  quadratic decision boundary
- What is  $C_{k,l}$ ? What if  $\boldsymbol{\Sigma}_k = \boldsymbol{\Sigma}_l$ ?

# Decision Boundary



# Learning

- Learn the parameters for each class using maximum likelihood
- Assume the prior is Bernoulli (we have two classes)

$$p(t|\phi) = \phi^t(1-\phi)^{1-t}.$$

- You can compute the MLE in closed form (good exercise!)

$$\phi = \frac{1}{N} \sum_{n=1}^N \mathbb{1}[t^{(n)} = 1]$$

$$\mu_k = \frac{\sum_{n=1}^N \mathbb{1}[t^{(n)} = k] \cdot \mathbf{x}^{(n)}}{\sum_{n=1}^N \mathbb{1}[t^{(n)} = k]}$$

$$\Sigma_k = \frac{1}{\sum_{n=1}^N \mathbb{1}[t^{(n)} = k]} \sum_{n=1}^N \mathbb{1}[t^{(n)} = k] (\mathbf{x}^{(n)} - \mu_{t^{(n)}}) (\mathbf{x}^{(n)} - \mu_{t^{(n)}})^T$$

# Simplifying the Model


What if  $\mathbf{x}$  is high-dimensional?

- For Gaussian Bayes Classifier, if input  $\mathbf{x}$  is high-dimensional, then covariance matrix has many parameters  $O(D^2)$
- Save some parameters by using a shared covariance for the classes, i.e.  $\Sigma_k = \Sigma_l$ .
- Any other idea you can think of? (next lecture)
- MLE in this case:

$$\Sigma = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}^{(n)} - \mu_{t^{(n)}})(\mathbf{x}^{(n)} - \mu_{t^{(n)}})^T$$

- Linear decision boundary (verify this mathematically!).

## Decision Boundary: Shared Variances (between Classes)



# Gaussian Discriminative Analysis vs Logistic Regression

- Binary classification: If you examine  $p(t = 1|\mathbf{x})$  under GDA and assume  $\Sigma_0 = \Sigma_1 = \Sigma$ , you will find that it looks like this:

$$p(t|\mathbf{x}, \phi, \mu_0, \mu_1, \Sigma) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x})}$$

where  $\mathbf{w}$  is an appropriate function of  $(\phi, \mu_0, \mu_1, \Sigma)$ ,  $\phi = p(t = 1)$ . You derived this in hw2.

- GDA is similar to logistic regression (LR), parameter estimates are computed differently.
- When should we prefer GDA to LR, and vice versa?

# Gaussian Discriminative Analysis vs Logistic Regression

- GDA is a generative model, LR is a discriminative model.
- GDA makes stronger modeling assumption: assumes class-conditional data is multivariate Gaussian
- If this is true, GDA is asymptotically efficient
- But LR is more robust, less sensitive to incorrect modeling assumptions (what loss is it optimizing?)
- Many class-conditional distributions lead to logistic classifier. (You saw an example in hw2)
- When these distributions are non-Gaussian (true almost always), LR usually beats GDA

## Generative models - Recap

- GDA has quadratic, LR has linear decision boundary
- With shared covariance, GDA is similar to logistic regression.
- Generative models:
  - ▶ Flexible models, easy to add/remove class.
  - ▶ Handle missing data naturally
  - ▶ More "natural" way to think about things, but usually doesn't work as well.
- Tries to solve a hard problem in order to solve an easy problem.

## Appendix: MLE for univariate Gaussian

$$0 = \frac{\partial \ell}{\partial \mu} = -\frac{1}{\sigma^2} \sum_{i=1}^N \mathbf{x}^{(i)} - \mu$$

$$\begin{aligned}0 &= \frac{\partial \ell}{\partial \sigma} = \frac{\partial}{\partial \sigma} \left[ \sum_{i=1}^N -\frac{1}{2} \log 2\pi - \log \sigma - \frac{1}{2\sigma^2} (\mathbf{x}^{(i)} - \mu)^2 \right] \\&= \sum_{i=1}^N -\frac{1}{2} \frac{\partial}{\partial \sigma} \log 2\pi - \frac{\partial}{\partial \sigma} \log \sigma - \frac{\partial}{\partial \sigma} \frac{1}{2\sigma} (\mathbf{x}^{(i)} - \mu)^2 \\&= \sum_{i=1}^N 0 - \frac{1}{\sigma} + \frac{1}{\sigma^3} (\mathbf{x}^{(i)} - \mu)^2 \\&= -\frac{N}{\sigma} + \frac{1}{\sigma^3} \sum_{i=1}^N (\mathbf{x}^{(i)} - \mu)^2\end{aligned}$$

$$\hat{\mu}_{\text{ML}} = \frac{1}{N} \sum_{i=1}^N \mathbf{x}^{(i)}$$

$$\hat{\sigma}_{\text{ML}} = \sqrt{\frac{1}{N} \sum_{i=1}^N (\mathbf{x}^{(i)} - \mu)^2}$$